

Weekly Report - A rough Profiling Attempt to Ethash

MSc Project
Runchao Han

December 15, 2017

1 Progress of the Last Week

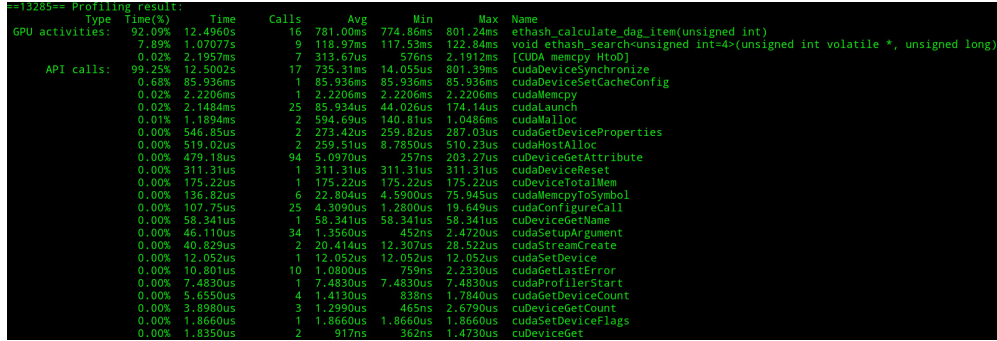
1. Did a rough profiling to Ethash by adding timestamps
2. Learned about Nvidia GPU and CUDA programming
3. Learned about profiling CUDA programs by `nvprof` and Nvidia Visual Profiler

2 Introduction to CUDA profilers

To profile CUDA programs, tools are involved. After searching for approaches, I found the only two profiling tools for CUDA programs:

- `nvprof`¹. This is a command line tool.
- Nvidia Visual Profiler². This is an Eclipse-based profiling tool which wraps `nvprof`.

A rough profiling was done by `nvprof`, by which a statistic was outputted, shown in Fig. 1.



Type	Time(%)	Time	Calls	Avg	Min	Max	Name
GPU activities:	92.09%	12.4960s	16	781.00ms	774.86ms	801.24ms	ethash_calculate_dag_item(unsigned int)
	7.89%	1.0707s	9	118.97ms	117.53ms	122.84ms	void ethash_search(unsigned int=4)(unsigned int volatile *, unsigned long)
	0.02%	2.1957ms	7	313.67us	576ms	2.1912ms	[CUDA memcpy HtoB]
API calls:	99.25%	12.5002s	17	735.31ms	14.05us	801.39ms	cudaDeviceSynchronize
	0.68%	85.936ms	1	85.936ms	85.936ms	85.936ms	cudaDeviceSetCacheConfig
	0.02%	2.2206ms	1	2.2206ms	2.2206ms	2.2206ms	cudaMemcpy
	0.02%	2.1484ms	25	85.934us	44.026us	174.14us	cudaLaunch
	0.01%	1.1894ms	2	594.69us	140.81us	1.0486ms	cudaMalloc
	0.00%	546.85us	2	273.42us	259.82us	287.03us	cudaGetDeviceProperties
	0.00%	519.02us	2	259.51us	8.7850us	510.23us	cudaHostAlloc
	0.00%	479.18us	94	5.0970us	257ms	203.27us	cudaDeviceGetAttribute
	0.00%	311.31us	1	311.31us	311.31us	311.31us	cudaDeviceReset
	0.00%	175.22us	1	175.22us	175.22us	175.22us	cudaDeviceTotalMem
	0.00%	136.82us	6	22.804us	4.5900us	75.945us	cudaMemcpyToSymbol
	0.00%	107.75us	25	4.3090us	1.2800us	19.649us	cudaConfigureCall
	0.00%	58.341us	1	58.341us	58.341us	58.341us	cudaDeviceGetName
	0.00%	46.110us	34	1.3560us	452ms	2.4720us	cudaSetupArgument
	0.00%	40.829us	2	20.414us	12.307us	28.522us	cudaStreamCreate
	0.00%	12.052us	1	12.052us	12.052us	12.052us	cudaSetDevice
	0.00%	10.801us	10	1.0800us	759ms	2.2330us	cudaGetLastError
	0.00%	7.4830us	1	7.4830us	7.4830us	7.4830us	cudaProfilerStart
	0.00%	5.6550us	4	1.4130us	838ms	1.7840us	cudaGetDeviceCount
	0.00%	3.8980us	3	1.2990us	465ms	2.6790us	cudaDeviceGetCount
	0.00%	1.8660us	1	1.8660us	1.8660us	1.8660us	cudaSetDeviceFlags
	0.00%	1.8350us	2	917ms	362ms	1.4730us	cudaDeviceGet

Figure 1: The `nvprof` result of `run_ethash_search()`.

The result shows that `ethash_calculate_dag_item()` costs most of the time, which generates the whole 1GB DAG by a seed. However, the DAG generation can be pre-computed or copied from

¹<http://docs.nvidia.com/cuda/profiler-users-guide/index.html>

²<http://www.sie.es/wp-content/uploads/2015/12/cuda-profiling-tools.pdf>

others, which makes optimising this process meaningless. The optimisation target is the mining process, which is the *compute_hash()* function in the code.

Currently, I have not succeeded in importing the project into **Nvidia Visual Profiler**. This is the main target next week.

3 Profiling *compute_hash()* Function

Due to the limitation of the command line, it is hard to get statistics of different steps which are set by myself. Therefore, I set timestamps manually in the code.

The steps of conducting the profiling are listed below:

1. Set the CUDA mining function involves only one block and the block contains only one thread.
 - A CUDA Kernel function takes a grid (fixed)
 - A grid takes several blocks (modifiable)
 - A block takes several threads (modifiable)
2. Add timestamps in the code. The code involves steps listed below:
 - (a) State initialisation
 - (b) An iteration of 4 (t)
 - i. An iteration of 4 (ti1)
 - ii. An iteration of 16 (ti2)
 - iii. An iteration of 4 (ti3)

The time of initialisation and three inner iterations of a single outer iteration is recorded by timestamps. (A problem is that if I record the whole outer iteration execution time, the long long int will even be overflowed. I will think about solving this next week.) The output is shown in Fig. 2.

The unit of values is the time of the GPU clock. The result indicates that the second inner iteration ti2 takes most of the time. However, it is uncertain that if different single outer iterations make inner iterations take different time. Further research is needed, which is the next week's plan.

It is noted that the implementation uses the most up-to-date CUDA APIs, like *__shfl()* and *__shfl_sync()*. Further research is needed to figure out these APIs.

4 Miscellaneous

- I decided to use Eclipse CDT to conduct experiments, which I found very convenient.
- I have learned basic CUDA programming by official documentations and examples. Currently I have basic understanding on the Ethminer code.

5 Next Week's Plan

1. Further profile the Ethash algorithm
2. Import the Ethminer to Nvidia Visual Profiler and get a better profiling result
3. Produce the Stack Graph if possible

No reference this week because the work in this week is about profiling the source code, which is fairly an engineering problem.

```
nonce-> 3128178979190234009;t1 -> 8652.000000
nonce-> 3128178979190234009;ti1 -> 1519.000000
nonce-> 3128178979190234009;ti2 -> 56217.000000
nonce-> 3128178979190234009;ti3 -> 680.000000

nonce-> 3128178979191282585;t1 -> 8878.000000
nonce-> 3128178979191282585;ti1 -> 1153.000000
nonce-> 3128178979191282585;ti2 -> 56133.000000
nonce-> 3128178979191282585;ti3 -> 633.000000

nonce-> 3128178979192331161;t1 -> 8903.000000
nonce-> 3128178979192331161;ti1 -> 1401.000000
nonce-> 3128178979192331161;ti2 -> 56043.000000
nonce-> 3128178979192331161;ti3 -> 743.000000

nonce-> 3128178979193379737;t1 -> 8654.000000
nonce-> 3128178979193379737;ti1 -> 1521.000000
nonce-> 3128178979193379737;ti2 -> 56320.000000
nonce-> 3128178979193379737;ti3 -> 681.000000

nonce-> 3128178979194428313;t1 -> 8661.000000
nonce-> 3128178979194428313;ti1 -> 1387.000000
nonce-> 3128178979194428313;ti2 -> 56803.000000
nonce-> 3128178979194428313;ti3 -> 684.000000

nonce-> 3128178979195476889;t1 -> 8817.000000
nonce-> 3128178979195476889;ti1 -> 1491.000000
nonce-> 3128178979195476889;ti2 -> 56320.000000
nonce-> 3128178979195476889;ti3 -> 938.000000

nonce-> 3128178979196525465;t1 -> 8652.000000
nonce-> 3128178979196525465;ti1 -> 1517.000000
nonce-> 3128178979196525465;ti2 -> 56451.000000
nonce-> 3128178979196525465;ti3 -> 771.000000
```

Figure 2: Profiling *compute_hash()* by adding timestamps.